

# NES Graphics

## CHR files

A NES ROM can be split into two files, whose extensions are .PRG and .CHR. Usually the PRG file is the content of the PRG ROM and the CHR file is the content of the CHR ROM, both inside the cartridge. The PRG chip connects to the NES motherboard's CPU and contains the compiled code programmed for the game, while the CHR chip goes to the PPU (Picture Processing Unit) and contains data for the graphics. The PCBs of the cartridges vary a lot in every game, so the aforementioned is not the rule. But to keep things simple and concise let's imagine that this is always the case.



Figure 1: PCB of a NES cartridge where we can see the CHR and PRG ROMS.

The sprites on the NES are 8x8 pixel *tiles*, and actually, the CHR contains the data for each of the tiles that are used on a particular game (and sometimes some that is not used) in a structured way. You can picture the contents of a CHR file as a matrix where the entries are 8x8 pixel tiles. Each tile only has 4 different colors, and usually many tiles are puzzled together to draw a complex shape in the screen.

The NES was released in 1983, technology back then was very limited compared to what we have nowadays. So it's not hard to realize that the size of a game was HUGE issue. The graphics had to be stored in an optimal way so data occupies the less space in memory as possible. The NES developers worked with hardware limitations in such a way that, for example, the size of the Super Mario Bros.nes file is 41KB. Imagine if we mapped all the stages of this game to computer images. If I take a screen-shot of my current screen on my computer, the file is about 20 times bigger than the size of the whole Super Mario Bros.nes file.

Why does this happen? It's reasonable that a screen-shot of my computer is much bigger than the whole game. My screen is displaying millions of colors, when the NES outputs only 64 colors (which there's 56 that are unique). Also, the NES based its screen output on the repetition of a definite amount of tiles, taking care that each of these tiles are stored only once in the game's data. Also, there's an important technique used to store all the graphic data that also reduced notoriously the amount of memory needed. This document's main goal is to describe this technique, which is basically indexed color on sub-palettes based on the 64 color NES palette. Let's imagine this situation: As  $2^6 = 64$ , to represent one of each of these 64 colors we would need 6 bits. So to represent a tile, we would need  $6 \times (8 \times 8) = 384$  bits. This is a lot of information and a waste of space, since we know that in the NES a tile can't have more than 4 different colors. To optimize this, every pixel is represented in the CHR chip with only 2 bits, so we only need 128 bits for each tile (that's  $\frac{1}{3}$  of what we needed before!). This means that if we assign a color to each of these  $2^2 = 4$  combinations, all the pixels in a CHR would have only 4 different colors. The way the NES has to decode the original colors for each tile is to assign a 'sub-palette' to each tile. A sub-palette has only 4 colors, taken from the 64 color NES palette, and these sub-palette / tile mapping are defined in the PRG chip.



Figure 2: NES palette.

**A quick example:** Imagine that the first 8 pixels of a tile are encoded with 01 00 00 00 11 11 11 10. Then suppose that the sub-palette mapped to this particular tile is 000110 110101 010101 000001, which represents a color from the 64 color palette. Then, with this mapping, the real colors of the first 8 pixels of the tile are: 110101 000110 000110 000110 000001 000001 000001 010101. This is, the first pixel is the color in the position 1 of the sub-palette: as positions are numbered from 0 to 3, that's 110101. The second, third and fourth pixels are the color in the position 0 of the sub-palette (000110). Fifth, sixth and seventh are the color in the position 3 of the sub-palette (000001). And the eight pixel is the color in the position 2 in the sub-palette (110101).

So basically each pixel is represented by an index for its corresponding palette. Knowing the corresponding palette would allow you to know the real color of the pixel, but for this you would need access to the game's code (not compiled).




Figure 3: This example gives you an idea of how we can represent data obtained from a CHR file. Here we can see 16 tiles where color 00 is black, 01 is red, 10 blue and 11 is white.




Figure 4: This shows in a graphical way a sub-palette mapped to the first tile from Figure 3. To the left we can see the palette and to the right the tile as we would see it on screen.

Now let's examine the contents of a CHR file. If you open one of these files you will find a large set of hexadecimal numbers. You can see these numbers are structured in a certain way. Remember that we earlier pictured the contents of a CHR file as matrix of tiles? The data inside these files is actually organized as huge matrix of 8 columns and  $n$  rows. The amount of rows varies and depends of the amount of graphics the game has. Each entry in this matrix consists of four hexadecimal digits.

A portion of the contents of a CHR file could be:

1	0000	0000	0000	0000	0000	0000	0000	0000
2	0000	0000	0000	0000	0000	0000	0000	0000
3	070f	3f7f	ffff	ff7f	0007	0f37	7f7f	6f1f
4	3f3f	1e1c	1c30	70fc	1e1c	0d0b	030f	2f73
5	60f0	f8f8	feff	ffff	0060	f0e0	f8e0	8000
6	c033	73f7	f7f7	3241	3fcc	be49	4949	ffbe
7	0000	0000	f0e0	d0f0	0000	0000	0000	0000
8	e040	8080	80c0	6070	0080	0000	0000	80a0

This is actually a small portion of the CHR file of the game Panic Restaurant - this game's CHR file is a 8192x8 matrix. In general, each row of these matrices represents one 8x8 pixel tile, encoded in a certain way. If we decode these matrices we are actually obtaining the all the uncolored tiles of a video-game. To understand the decoding

process, we'll go through an example. Let's decode the values for each pixel of the tile from the row six of the excerpt of Panic Restaurant's matrix: **c033 73f7 f7f7 3241 3fcc be49 4949 ffbe**.

First, let's split the original row in two parts and place them like this:

```
1 c033 73f7 f7f7 3241
2 3fcc be49 4949 ffbe
```

Then, let's obtain the binary representation for each entry of the row:

```
1 0xC033 = 1100000000110011b
2 0x73F7 = 0111001111110111b
3 0xF7F7 = 1111011111110111b
4 0x3241 = 0011001001000001b
5 0x3FCC = 0011111111001100b
6 0xBE49 = 1011111001001001b
7 0x4949 = 0100100101001001b
8 0xFFBE = 1111111110111110b
```

And replace the split row with its binary representation:

```
1 1100 0000 0011 0011 0111 0011 1111 0111 1111 0111 1111 0111 0011 0010 0100 0001
2 0011 1111 1100 1100 1011 1110 0100 1001 0100 1001 0100 1001 1111 1111 1011 1110
```

Notice that here we had added a white-space between each hex digits binary representation for a better appreciation. Also we can see that we have 128 bits of data, which is the amount of bits we calculated before that would be needed for a tile.

So in each part of the split row we have 64 bits. As each tile is  $8 \times 8 = 64$  pixels, this suggests there is a relation between each part of the split row. The relation is very simple and it's based on the position of each bit. From now on, let's think about the two parts of the split row as line one and line two.

Imagine each of these lines as two vectors  $\phi_1$  and  $\phi_2$ . Each bit for each line (excluding the white-spaces we added on purpose) is an element on its corresponding vector. In fact let's force that the elements of the vectors appear in the same order as they are on the lines. If any doubt: the value of the element in the position  $k$  in the vector is the bit in the  $k$  position of its corresponding line (again, excluding white-spaces). Also let's define an access operation/function:

$$\phi_i[j], (i \in \{1, 2\}; j \in \{1, \dots, 64\})$$

that retrieves the decimal number that represents the value of the entry at the  $j$  position in the vector  $\phi_i$ .

Now, let's define a new vector  $\psi$  of length 64, where:

$$\psi[j] = (\phi_1[j] \times 1) + (\phi_2[j] \times 2)$$

We have applied a decoding function taking corresponding elements from  $\phi_1$  and  $\phi_2$  and built a new vector that actually is the representation of a tile. Then if we decompose this vector into an  $8 \times 8$  matrix we can see the 'shape' of the tile. The function applied to the two vectors is very simple, for each of the 64 possible positions, take the corresponding value from the first vector and add two times the value of the corresponding position in the second vector. The result is the position of the color of the pixel on the corresponding sub-palette. There's two minor things I would like to mention here. First is that  $\psi$  is a vector that contains decimal numbers. As we defined the access operation as (some sort of) a function that retrieves a decimal value, the  $+$  and  $\times$  operations that appear in the decoding function are the base-10 addition and multiplication operators. This is just to make things look

nicer. Second thing is that all the possible resulting values of applying this function are 0, 1, 2, 3, so  $\psi[j] \in \{0, 1, 2, 3\}$ ,  $j \in \{1, \dots, 64\}$ .

For example, the 4th digit of  $\phi_1$  is 0, and the 4th digit of  $\phi_2$  is 1, so the result is:

$$(\phi_1[3] \times 1) + (\phi_2[3] \times 2) = (0 \times 1) + (1 \times 2) = 2$$

This means the 4th pixel's color is the one in the position  $2_{10}$  (or  $10_2$ ) of the tile's corresponding palette.

We could calculate the codes for each pixel of the tile by hand pretty easy, but it would be a waste of time. Let's do some **Python** code to see some quick results:

```
1 line = 'c033_73f7_f7f7_3241_3fcc_be49_4949_ffbe'.replace('_', '')
2
3 result = \
4 [ x + 2*y for x, y in \
5     zip(\
6         [int(x) for x in list(bin(int(line[:16], 16))[2:].zfill(64))], \
7         [int(x) for x in list(bin(int(line[16:], 16))[2:].zfill(64))]\
8     )\
9 ]
```

There's a much simpler, explained equivalent version of this code at the end of the document.

This code outputs the following:

[1, 1, 2, 2, 2, 2, 2, 2, 2, 1, 1, 2, 2, 1, 1, 2, 1, 3, 3, 2, 2, 3, 1, 1, 3, 1, 1, 2, 1, 1, 3, 1, 3, 1, 1, 2, 1, 1, 3, 1, 3, 1, 1, 2, 1, 1, 3, 2, 2, 3, 3, 2, 2, 3, 2, 2, 1, 2, 2, 2, 2, 2, 1]

Which is a position on the sub-palette for each of the 64 pixels of the tile. Let's organize this in an 8x8 matrix:

```
1 matrix = \
2 [ result[index:index + 8] for index in range(0, len(result), 8) ]
3
4 print(*matrix, sep='\n')
```

That code outputs:

```
[1, 1, 2, 2, 2, 2, 2, 2]
[2, 2, 1, 1, 2, 2, 1, 1]
[2, 1, 3, 3, 2, 2, 3, 1]
[1, 3, 1, 1, 2, 1, 1, 3]
[1, 3, 1, 1, 2, 1, 1, 3]
[1, 3, 1, 1, 2, 1, 1, 3]
[2, 2, 3, 3, 2, 2, 3, 2]
[2, 1, 2, 2, 2, 2, 2, 1]
```

It would be really easy to get the binary representation of each entry if we need to:

```
1 print(*[ [ bin(entry)[2:].zfill(2) for entry in result ][index:index + 8] \
2         for index in range(0, len(result), 8) ], \
3         sep='\n'\
4     )
```

The output is:

```
['01', '01', '10', '10', '10', '10', '10', '10']
['10', '10', '01', '01', '10', '10', '01', '01']
['10', '01', '11', '11', '10', '10', '11', '01']
['01', '11', '01', '01', '10', '01', '01', '11']
['01', '11', '01', '01', '10', '01', '01', '11']
['01', '11', '01', '01', '10', '01', '01', '11']
['10', '10', '11', '11', '10', '10', '11', '10']
['10', '01', '10', '10', '10', '10', '10', '01']
```

Note that in this example there's no pixel with the color of position 0 of the sub-palette.

Let's calculate the colors for each row/tile on the portion of the CHR file originally listed:

```
1 tiles = ['0000_0000_0000_0000_0000_0000_0000_0000', \
2         '0000_0000_0000_0000_0000_0000_0000_0000', \
3         '070f_3f7f_ffff_ff7f_0007_0f37_7f7f_6f1f', \
4         '3f3f_1e1c_1c30_70fc_1e1c_0d0b_030f_2f73', \
5         '60f0_f8f8_ffff_ffff_0060_f0e0_f8e0_8000', \
6         'c033_73f7_f7f7_3241_3fcc_be49_4949_ffbe', \
7         '0000_0000_f0e0_d0f0_0000_0000_0000_0000', \
8         'e040_8080_80c0_6070_0080_0000_0000_80a0'
9     ]
10
11 for tile in tiles:
12     line = tile.replace('_', ' ')
13     result = \
14         [ x + 2*y for x, y in \
15             zip(\
16                 [int(x) for x in list(bin(int(line[:16], 16))[2:].zfill(64))], \
17                 [int(x) for x in list(bin(int(line[16:], 16))[2:].zfill(64))]\
18             )\
19         ]
20     matrix = \
21         [ result[index:index + 8] for index in range(0, len(result), 8) ]
22     print(*matrix, sep='\n')
23     print('\n')
```

tile 1:	tile 2:	tile 3:	tile 4:	tile 5:
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 1, 1, 1]	[0, 0, 1, 3, 3, 3, 3, 1]	[0, 1, 1, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 1, 3, 3, 3]	[0, 0, 1, 3, 3, 3, 1, 1]	[1, 3, 3, 1, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 1, 1, 3, 3, 3, 3]	[0, 0, 0, 1, 3, 3, 1, 2]	[3, 3, 3, 3, 1, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[0, 1, 3, 3, 1, 3, 3, 3]	[0, 0, 0, 1, 3, 1, 2, 2]	[3, 3, 3, 1, 1, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 3, 3, 3, 3, 3, 3, 3]	[0, 0, 0, 1, 1, 1, 2, 2]	[3, 3, 3, 3, 3, 1, 1, 0]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 3, 3, 3, 3, 3, 3, 3]	[0, 0, 1, 1, 2, 2, 2, 2]	[3, 3, 3, 1, 1, 1, 1, 1]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 3, 3, 1, 3, 3, 3, 3]	[0, 1, 3, 1, 2, 2, 2, 2]	[3, 1, 1, 1, 1, 1, 1, 1]
[0, 0, 0, 0, 0, 0, 0, 0]	[0, 0, 0, 0, 0, 0, 0, 0]	[0, 1, 1, 3, 3, 3, 3, 3]	[1, 3, 3, 3, 1, 1, 2, 2]	[1, 1, 1, 1, 1, 1, 1, 1]

tile 6:	tile 7:	tile 8:
[1, 1, 2, 2, 2, 2, 2, 2]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 1, 1, 0, 0, 0, 0, 0]
[2, 2, 1, 1, 2, 2, 1, 1]	[0, 0, 0, 0, 0, 0, 0, 0]	[2, 1, 0, 0, 0, 0, 0, 0]
[2, 1, 3, 3, 2, 2, 3, 1]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 0, 0, 0, 0, 0, 0, 0]
[1, 3, 1, 1, 2, 1, 1, 3]	[0, 0, 0, 0, 0, 0, 0, 0]	[1, 0, 0, 0, 0, 0, 0, 0]
[1, 3, 1, 1, 2, 1, 1, 3]	[1, 1, 1, 1, 0, 0, 0, 0]	[1, 0, 0, 0, 0, 0, 0, 0]
[1, 3, 1, 1, 2, 1, 1, 3]	[1, 1, 1, 0, 0, 0, 0, 0]	[1, 1, 0, 0, 0, 0, 0, 0]
[2, 2, 3, 3, 2, 2, 3, 2]	[1, 1, 0, 1, 0, 0, 0, 0]	[2, 1, 1, 0, 0, 0, 0, 0]
[2, 1, 2, 2, 2, 2, 2, 1]	[1, 1, 1, 1, 0, 0, 0, 0]	[2, 1, 3, 1, 0, 0, 0, 0]

To conclude, by understanding this we know how the graphic data is stored inside the CHR ROM of the NES cartridges and also how the NES works with these graphics. With some work, we could be able to convert our own graphics to NES format. Also, we can read the contents of a CHR file and calculate the position on the sub-palette of each pixel. Note that most of the tiles are placed next to other tiles to create compound, bigger tiles... So if we assign a random color to each different number on a decoded CHR file (4 different numbers = 4 different colors) still many of these tiles would draw reasonable shapes even without their original colors. A nice experiment would be to read the contents of a whole CHR file and process the data. Then we can find tiles that go together and re-build some graphics as characters, items, etc.

# Addendum

## Resources

Some useful tools to check out what I wrote in this document.

### Sublime Text Editor

This is a nice cross-platform text editor that can open CHR files and display its contents. I've tried other editors such as Atom and Xed and it didn't work well.

### FamiRom

This is a nice tool that lets you load a NES ROM file and splits it into PRG and CHR files. The software is for Windows, but you can use it on Linux without problems if you install the package **mono-complete**. Then just run **mono famiRom.exe** in the linux console.

### Python

A very powerful object oriented programming language.

## Python code for dummies!

```
1 # Create two strings with each line of hexadecimal values and removing whitespaces
2 l1 = 'c033_73f7_f7f7_3241'.replace('_', '')
3 l2 = '3fcc_be49_4949_ffbe'.replace('_', '')
4
5 # Define a function that given a string representation of a hexadecimal number
6 # returns a binary string of that number, removing the first two characters that
7 # are '0b' for binary. Fills with zeros to complete 4 digits.
8 def to_bin(hexa_num):
9     to_int = int(hexa_num, 16)
10    return bin(to_int)[2:].zfill(4)
11
12 # The function to calculate the color
13 def calculate(num1, num2):
14    return num1 + (2 * num2)
15
16 # Creating a string of binary digits
17 # for each line of hexadecimal numbers
18 bin_string1 = ''
19 bin_string2 = ''
20
21 for index in range(0, 16):
22     bin_string1 += (to_bin(l1[index]))
23     bin_string2 += (to_bin(l2[index]))
24
25 # Turns the strings to lists, then for each entry of the list, converts to
26 # integer and apply the function to calculate the resulting color.
27 list1 = list(bin_string1)
28 list2 = list(bin_string2)
29 res = list()
30
31 for index in range(0, 64):
32     num1 = int(list1[index])
33     num2 = int(list2[index])
34     res.append(calculate(num1, num2))
```